

RECOMMENDATION

Reactivate the dormant pipeline on registry signals, not on a cadence

15,000+ clinicians have tried JotPsych. Thousands pay nothing today, and most never said no — they were mid-contract or mid-EHR-decision. That pool is already-paid-for pipeline: acquisition cost was spent once, and at current pricing each reactivation is worth **\$637–\$3,229/yr** (Moonlighting to Everything, annual billing). Second Window works that pool with a machine instead of a sequence: it matches each name to the federal NPI registry, writes only when a practice visibly changes shape — the moments an EHR decision reopens — and stays silent otherwise. It is deployed and running.

Why this is worth running

Revenue math	Every 1% of a 4,000-name dormant pool that returns at the Basic tier is ≈ \$65K ARR . At Everything, ≈ \$129K. The machine exists to find which 1–3% are reachable and when.
Channel cost	≈ \$10–20/month all-in (one small VM, model spend at ~\$0.02 per drafted message, registry API is free, email free tier). One Moonlighting reactivation pays for roughly three years of running it.
Second market, free	151 behavioural-health practices registered with CMS in the last 90 days, found by one capped sweep of 10 states — practices choosing their first EHR, before any competitor is in the room. Zero list cost, zero acquisition spend.
Labour	1–2 hours/month, fixed, independent of list size: ≤10 pre-written peer introductions per cycle. Everything else is autopilot behind checks strict enough that they currently block first and ask questions never — 11 drafts stopped on the live deployment to date.

What the machine reports, and the decisions those numbers drive

Every cycle logs, per trigger type (practice moved, went solo-to-group, changed specialty, newly registered): messages sent, replies, demos, returns, and **touches per return**. Silence is logged too — 93% of decisions on the sample audience — so restraint is measured, not claimed. Pre-committed decision rules: after 90 days on the real list, **keep** any trigger converting to a demo at ≥1%, **cut** any angle with 200+ sends and zero returns, and **move** the confidence thresholds only on the logged precision-by-band curve. If no trigger converts, the detector thesis is wrong and the numbers will say so within one quarter — that is the cheapest possible way to find out.

Current state (live deployment)

39 clinicians watched · 151 new practices surfaced · 5 cycles · 11 drafts blocked by checks · 1 return recorded (0 attributed — attribution is computed from the machine's own send history; it takes no credit for returns it did not touch). Sending to real clinicians is off pending a go decision; every other part runs unattended, including return capture via webhook from the signup flow.

Next 30 days

1. Load the real dormant list (one CSV upload or a Sheets/CRM URL; the machine needs only name, email, mobile).
2. Point the signup webhook at production so returns attribute automatically.
3. Switch to the CMS bulk files — the weekly incremental is a ready-made change feed; 15,000 lookups become one download.
4. Turn sending on at 3/day behind the existing caps, read the first fortnight of numbers, then scale what converts.

Live: jotpsych-second-window.fly.dev (open; sending off; sample messages deliver on request) · Code: github.com/joshxnyc/JotPsychCOS · Sample data labeled throughout; registry records are real public NPPES data. Proof it runs, including what it refused to send, follows on pages 2–4.

PROOF

One full pass, start to finish

Reproduced from a real run in the repository. Cycle 1 reads a list the machine did not write, matches every clinician against the live federal registry, and builds its baseline. The seed tool then rolls five records back to a plausible earlier state — the one simulated input, labeled, because a first run has no “before” to diff against. Cycle 2 sees the changes through the ordinary code path: one message staged, one clinician handed to a person, 35 deliberately left alone.

```
$ ./run.sh --limit 5 # cycle 1: builds the baseline
[input] 40 rows from inbox/dormant.sample.csv [SAMPLE DATA]
[input] 9 rows from inbox/peers.sample.csv [SAMPLE DATA]
[input] 3 rows from inbox/returns.sample.csv [SAMPLE DATA]
[input] 2 rows from inbox/suppress.sample.csv [SAMPLE DATA]
[input] 40 dormant clinicians, 9 peers, 3 return events
[resolve] re-read 38 of 38 from the registry (budget 400); 0 reused from memory
[decide] {'silence': 34, 'keep_warm': 3}
[stage] outbox(simulated) -> sarahgoldberg@icloud.com :: Running JotPsych beside what you already use
[stage] outbox(simulated) -> odalys@fkbehavioralhealth.com :: Running JotPsych beside what you already use
[stage] outbox(simulated) -> a.sorokin@lakeshoremindhealth.com :: Running JotPsych beside what you already use
[digest] outbox(simulated) -> operator@jotpsych.com

[done] staged=3 blocked=0 silent=34 human_queue=0 skipped=0
[done] report /home/user/JotPsychCOS/out/index.html
[done] queue /home/user/JotPsychCOS/out/human_queue.md

$ python tools/seed_prior_snapshot.py # simulated 'before' (labeled)
(no record suitable for a entity rollback – skipped)
marked 5 profiles stale so the next run re-reads them
Rolled back 5 of 20 records to a simulated prior state:
0d99fce07bda state: 'NJ' -> (next run will see) 'CA'
15c5c1861afb city: 'BROOKLYN' -> (next run will see) 'GARDENA'
19ef84aecb5c taxonomy: 'Counselor, Professional' -> (next run will see) 'Social Worker, Clinical'
3a8f99227848 address_1: '100 MAIN ST STE 2' -> (next run will see) '6 OLDE HICKORY PATH'
522f104b8680 name: 'MICHAEL SOLO PRACTICE' -> (next run will see) 'MICHAEL CHEN'

SIMULATED. Delete state/registry_snapshot.json to start clean.

$ ./run.sh --limit 6 # cycle 2: the diff has something to see
[input] 40 rows from inbox/dormant.sample.csv [SAMPLE DATA]
[input] 9 rows from inbox/peers.sample.csv [SAMPLE DATA]
[input] 3 rows from inbox/returns.sample.csv [SAMPLE DATA]
[input] 2 rows from inbox/suppress.sample.csv [SAMPLE DATA]
[input] 40 dormant clinicians, 9 peers, 3 return events
[resolve] re-read 6 of 38 from the registry (budget 400); 32 reused from memory
[decide] {'human_call': 1, 'moment': 1, 'silence': 35}
[stage] outbox(simulated) -> jnguyen@orangegrovetherapy.com :: Running JotPsych beside what you already use
[digest] outbox(simulated) -> operator@jotpsych.com

[done] staged=1 blocked=0 silent=35 human_queue=1 skipped=0
[done] report /home/user/JotPsychCOS/out/index.html
[done] queue /home/user/JotPsychCOS/out/human_queue.md
```

An output that left the program

A staged message as a real RFC-822 email file — the same code path as a live send, one switch away, with the simulation declared in its headers and the compliance footer (hyperlinked unsubscribe in the HTML part, bare URL in text, postal address in both) appended by the machine, not the model.

```
From: JotPsych Machine <onboarding@resend.dev>
To: a.sorokin@lakeshoremindhealth.com
Subject: Running JotPsych beside what you already use
X-Machine-Simulated: STAGED - passed every check, awaiting approval

You looked at JotPsych a while back and stayed where you were. That is usually the right call
mid-contract.

Worth knowing: you do not have to move anything. JotPsych runs alongside the system you are on today,
and it takes about five minutes to set up. Notes and claims get checked against payer rules before they
go out, which is where most denials start.

If it is useful later, it will still be here.

- Josh, JotPsych

--
You are receiving this because you signed up for JotPsych.
Unsubscribe: https://jotpsych-second-window.fly.dev/unsubscribe?t=a.sorokin@lakeshoremindhealth.com:bb0f283705394d29
JotPsych, Brooklyn Navy Yard, Dock 72, 7th Floor, Brooklyn, NY 11205
```

The hour of human work it produced

Generated by cycle 2. Not “review the output” — a specific introduction with the opening line written.

```
# The human hour

_Generated 2026-08-18 20:57 UTC · run 2_

These are the only clinicians this cycle where a person beats a message.
Each one hit the strongest class of signal **and** has a consenting peer
in their specialty and state. Twelve minutes each. Nothing else needs you.

## 1. Jordan Blake
- **Why now:** their practice moved to another state
- **Identity confidence:** probable (57)
- **Peer to offer:** Tessa Nolan, LMHC — Behavior Technician in CA, 4 months on JotPsych
- **Open with:** "Supervision notes and session notes finally live in the same place."
- **Do not say:** that anything was looked up. You are offering an introduction, not reporting on their
practice.

---

**Also worth your time this month:**
- Skim `out/quarantine/` — if the machine is blocking the same thing repeatedly, the fix is a line in
`config/guardrails.yaml`, not a rewrite.
- Check the angle table. Retire anything with sends and no returns after two cycles.
```

What it refused to send, and the check that caught it

The most important refusal this machine makes. Using the public registry to choose *when* to write is legitimate; telling a clinician they were looked up is surveillance, and a clinician who feels watched is lost for good. This draft was pushed through the unmodified production checks by the adversarial harness (tools/red_team.py) and quarantined:

The draft, as written

```
To: clinician@example.com
Subject: Congratulations on the new practice

I noticed you recently registered a new practice location. JotPsych runs alongside the system you
already use and takes about five minutes to set up. Notes and claims are checked against payer rules
before they go out. Reply stop and I will take you off the list. JotPsych runs alongside the system you
already use and takes about five minutes to set up. Notes and claims are checked against payer rules
before they go out. Reply stop and I will take you off the list.

--
You are receiving this because you signed up for JotPsych.
To stop receiving these, reply with the word STOP.
JotPsych, Brookly
```

The verdict

```
BLOCKED — reveals how we knew: 'i noticed'
quarantined as out/quarantine/20260818T205757-clinician-example-com-a91f2f.json
```

Every check, exercised on purpose

Eight drafts through the same code path: seven must be blocked, and one clean draft must pass — so the gates cannot quietly degrade into blocking everything.

```
[PASS] says how we knew
[PASS] claims a fact the identity confidence has not earned
[PASS] guarantees a reimbursement outcome
[PASS] invents a compliance credential
[PASS] reads as unread AI output
[PASS] echoes the registry change back at them
[PASS] leaks something that looks like patient data
[PASS] a clean draft, which must pass
All 8 behaved correctly. 7 blocked and quarantined, 1 clean draft allowed through.
```

And once with nobody watching

During a scheduled run in CI, the machine's judge rejected three drafts written by the real model — objecting that “You tried JotPsych once” was a claim about the recipient the record did not support. It was right: the fact was true of everyone on the list, but the prompt had never said so, and the judge refused to send an inference. The bug was in the prompt, not the model — caught by an unsupervised check, which is the property the whole design is built around: when the judge cannot return a verdict, or the draft overreaches, the message is quarantined, never sent. The live deployment shows the same behaviour today: 11 drafts stopped, each with its reason visible in the console.