RECOMMENDATION

Build the readiness detector, not the drip

Thousands of clinicians tried JotPsych and did not subscribe. Most did not say no — the timing was wrong: mid-contract, mid-EHR-decision, not ready. A nurture sequence answers that with volume, and volume is how a brand gets tuned out by the exact audience it needs back. So **Second Window** uses the three fields a signup leaves behind (name, email, mobile) to match each clinician to the federal NPI registry, then watches that registry for the practice changes that reopen a software decision — a solo NPI becoming a group, a move, a new specialty, a practice being formed. It writes at that moment, checks every word against what JotPsych can honestly claim, and stays silent the rest of the time. Silence is its most common output, and it is measured.

The numbers it reports

93% silence rate	of decisions on the sample audience were a deliberate 'say nothing', each logged with its reason. Restraint at 15,000 names is the product.
39 watched · 151 new practices	live right now: the uploaded audience, plus behavioural-health NPIs registered with CMS in the last 90 days — practices choosing their first system, found with no list from anyone.
7 drafts stopped	by its own checks on the live deployment — including three written by the real model and rejected by the machine's own judge, unsupervised.
1 return recorded, 0 attributed	attribution is computed from the machine's own send history, never asserted: “the machine never wrote to them, so it takes no credit.”

The one to two hours a month

Not spent reading output — the machine runs on autopilot, and a message that passes every check sends on its own. The hours go to the one thing software cannot do: at most ten warm introductions a cycle, each pairing a detected moment with a consenting peer in the same specialty and state, the email already written and checked, the peer's own words quoted verbatim. Twelve minutes each. The monthly remainder: skim what the checks stopped, and retire any angle with sends and no returns.

Week two

1. Swap per-clinician registry lookups for the CMS bulk files — the weekly incremental file is literally a change feed, turning 15,000 API calls into one download and making detection same-week and complete rather than sampled. **2.** Point the signup webhook (already built) at JotPsych's real signup flow, so returns attribute themselves with nobody in the loop. **3.** Calibrate the confidence thresholds and trigger weights on outcomes — the machine already logs precision by trigger type, so if practice moves convert and specialty changes do not, that is a settings change, not a rewrite. The open question it is built to answer: whether a registry change predicts buying. If none of the triggers convert, I would rather learn that in week three than year two.

Live workspace: **jotpsych-second-window.fly.dev** (open, no sign-up — sending to clinicians is off; sample messages deliver to whoever asks). Code and full documentation: **github.com/joshxnyc/JotPsychCOS**. Sample data is labeled as such throughout; registry records are real public NPPES data.

PROOF

One full pass, start to finish

Reproduced from a real run in the repository. Cycle 1 reads a list the machine did not write, matches every clinician against the live federal registry, and builds its baseline. The seed tool then rolls five records back to a plausible earlier state — the one simulated input, labeled, because a first run has no “before” to diff against. Cycle 2 sees the changes through the ordinary code path: one message staged, one clinician handed to a person, 35 deliberately left alone.

```
$ ./run.sh --limit 5 # cycle 1: builds the baseline
[input] 40 rows from inbox/dormant.sample.csv [SAMPLE DATA]
[input] 9 rows from inbox/peers.sample.csv [SAMPLE DATA]
[input] 3 rows from inbox/returns.sample.csv [SAMPLE DATA]
[input] 2 rows from inbox/suppress.sample.csv [SAMPLE DATA]
[input] 40 dormant clinicians, 9 peers, 3 return events
[resolve] re-read 38 of 38 from the registry (budget 400); 0 reused from memory
[decide] {'silence': 34, 'keep_warm': 3}
[stage] outbox(simulated) -> sarahgoldberg@icloud.com :: Running JotPsych beside what you already use
[stage] outbox(simulated) -> odalys@fkbehavioralhealth.com :: Running JotPsych beside what you already use
[stage] outbox(simulated) -> a.sorokin@lakeshoremindhealth.com :: Running JotPsych beside what you already use
[digest] outbox(simulated) -> operator@jotpsych.com

[done] staged=3 blocked=0 silent=34 human_queue=0 skipped=0
[done] report /home/user/JotPsychCOS/out/index.html
[done] queue /home/user/JotPsychCOS/out/human_queue.md

$ python tools/seed_prior_snapshot.py # simulated 'before' (labeled)
(no record suitable for a entity rollback — skipped)
marked 5 profiles stale so the next run re-reads them
Rolled back 5 of 20 records to a simulated prior state:
0d99fce07bda state: 'NJ' -> (next run will see) 'CA'
15c5c1861afb city: 'BROOKLYN' -> (next run will see) 'GARDENA'
19ef84aecb5c taxonomy: 'Counselor, Professional' -> (next run will see) 'Social Worker, Clinical'
3a8f99227848 address_1: '100 MAIN ST STE 2' -> (next run will see) '6 OLDE HICKORY PATH'
522f104b8680 name: 'MICHAEL SOLO PRACTICE' -> (next run will see) 'MICHAEL CHEN'

SIMULATED. Delete state/registry_snapshot.json to start clean.

$ ./run.sh --limit 6 # cycle 2: the diff has something to see
[input] 40 rows from inbox/dormant.sample.csv [SAMPLE DATA]
[input] 9 rows from inbox/peers.sample.csv [SAMPLE DATA]
[input] 3 rows from inbox/returns.sample.csv [SAMPLE DATA]
[input] 2 rows from inbox/suppress.sample.csv [SAMPLE DATA]
[input] 40 dormant clinicians, 9 peers, 3 return events
[resolve] re-read 6 of 38 from the registry (budget 400); 32 reused from memory
[decide] {'human_call': 1, 'moment': 1, 'silence': 35}
[stage] outbox(simulated) -> jnguyen@orangegrovetherapy.com :: Running JotPsych beside what you already use
[digest] outbox(simulated) -> operator@jotpsych.com

[done] staged=1 blocked=0 silent=35 human_queue=1 skipped=0
[done] report /home/user/JotPsychCOS/out/index.html
[done] queue /home/user/JotPsychCOS/out/human_queue.md
```

An output that left the program

A staged message as a real RFC-822 email file — the same code path as a live send, one switch away, with the simulation declared in its headers and the compliance footer (hyperlinked unsubscribe in the HTML part, bare URL in text, postal address in both) appended by the machine, not the model.

```
From: JotPsych Machine <onboarding@resend.dev>
To: a.sorokin@lakeshoremindhealth.com
Subject: Running JotPsych beside what you already use
X-Machine-Simulated: STAGED - passed every check, awaiting approval

You looked at JotPsych a while back and stayed where you were. That is usually the right call
mid-contract.

Worth knowing: you do not have to move anything. JotPsych runs alongside the system you are on today,
and it takes about five minutes to set up. Notes and claims get checked against payer rules before they
go out, which is where most denials start.

If it is useful later, it will still be here.

— Josh, JotPsych

--
You are receiving this because you signed up for JotPsych.
Unsubscribe: https://jotpsych-second-window.fly.dev/unsubscribe?t=a.sorokin@lakeshoremindhealth.com:bb
0f283705394d29
JotPsych, Brooklyn Navy Yard, Dock 72, 7th Floor, Brooklyn, NY 11205
```

The hour of human work it produced

Generated by cycle 2. Not “review the output” — a specific introduction with the opening line written.

```
# The human hour

_Generated 2026-08-18 20:57 UTC · run 2_

These are the only clinicians this cycle where a person beats a message.
Each one hit the strongest class of signal **and** has a consenting peer
in their specialty and state. Twelve minutes each. Nothing else needs you.

## 1. Jordan Blake
- **Why now:** their practice moved to another state
- **Identity confidence:** probable (57)
- **Peer to offer:** Tessa Nolan, LMHC — Behavior Technician in CA, 4 months on JotPsych
- **Open with:** "Supervision notes and session notes finally live in the same place."
- **Do not say:** that anything was looked up. You are offering an introduction, not reporting on their
practice.

---

**Also worth your time this month:**
- Skim `out/quarantine/` — if the machine is blocking the same thing repeatedly, the fix is a line in
`config/guardrails.yaml`, not a rewrite.
- Check the angle table. Retire anything with sends and no returns after two cycles.
```

What it refused to send, and the check that caught it

The most important refusal this machine makes. Using the public registry to choose *when* to write is legitimate; telling a clinician they were looked up is surveillance, and a clinician who feels watched is lost for good. This draft was pushed through the unmodified production checks by the adversarial harness (tools/red_team.py) and quarantined:

The draft, as written

```
To: clinician@example.com
Subject: Congratulations on the new practice

I noticed you recently registered a new practice location. JotPsych runs alongside the system you
already use and takes about five minutes to set up. Notes and claims are checked against payer rules
before they go out. Reply stop and I will take you off the list. JotPsych runs alongside the system you
already use and takes about five minutes to set up. Notes and claims are checked against payer rules
before they go out. Reply stop and I will take you off the list.


--
You are receiving this because you signed up for JotPsych.
To stop receiving these, reply with the word STOP.
JotPsych, Brookly
```

The verdict

```
BLOCKED — reveals how we knew: 'i noticed'
quarantined as out/quarantine/20260818T205757-clinician-example-com-a91f2f.json
```

Every check, exercised on purpose

Eight drafts through the same code path: seven must be blocked, and one clean draft must pass — so the gates cannot quietly degrade into blocking everything.

```
[PASS] says how we knew
[PASS] claims a fact the identity confidence has not earned
[PASS] guarantees a reimbursement outcome
[PASS] invents a compliance credential
[PASS] reads as unread AI output
[PASS] echoes the registry change back at them
[PASS] leaks something that looks like patient data
[PASS] a clean draft, which must pass
All 8 behaved correctly. 7 blocked and quarantined, 1 clean draft allowed through.
```

And once with nobody watching

During a scheduled run in CI, the machine's judge rejected three drafts written by the real model — objecting that “You tried JotPsych once” was a claim about the recipient the record did not support. It was right: the fact was true of everyone on the list, but the prompt had never said so, and the judge refused to send an inference. The bug was in the prompt, not the model — caught by an unsupervised check, which is the property the whole design is built around: when the judge cannot return a verdict, or the draft overreaches, the message is quarantined, never sent. The live deployment shows the same behaviour today: 7 drafts stopped, each with its reason visible in the console.